HIERARCHICAL ACTIVE LEARNING (HAL) APPLICATION TO MITOCHONDRIAL DISEASE PROTEIN DATASET

James Duin

University of Nebraska–Lincoln
Master's Thesis

Spring 2017

jamesdduin@gmail.com

Introduction

HAL - Protein

James Duin

Introduction
Background
Exp. Setup
Conv. ML
Act. vs Pass.
HAL Results
BANDIT Results
Conclusions

- Identify the source of mutations which give rise to mitochondrial disease
- Leigh Syndrome, Lebers Hereditary Optic Neuropathy
- Hierarchically labeled according to location in mitochondria
- Coarse-grained: learning labels near the root of the tree
- Fine-grained: learning labels towards the leaf nodes
- Learn mitochondrion concept (coarse) by combining classifiers for each target compartment (fine)

Introduction
Related Work

HAL - Protein
James Duin
Introduction
Background
Exp. Setup
Conv. ML
Act. vs Pass.
HAL Results
BANDIT Results
Conclusions

- **Active learning**: copious unlabeled data, cost associated with acquiring labels, yields best classifier for a given cost, or best for minimal cost
- Previous work in text classification and and rich media indexing use hierarchies of labels to improve fine-level classification (McCallum et al. 1998, Jiang et al. 2013)
- Previous work in named entity recognition to target fine-grained entity categories (Fleischman et al. 2002)
- **First, investigation of active learning in a hierarchical setting, approach shown to find best classifier for a budget regardless of varying label acquisition cost**

Introduction

Outline:

- Machine Learning
- Active Machine Learning
- Coarse-grained vs Fine-grained Trade Off
- Active over-labeling algorithms
- Hierarchical Protein Dataset
- Application to Protein Dataset
- Experimental Results

HAL - Protein
James Duin
Introduction
Background
Exp. Setup
Conv. ML
Act. vs Pass.
HAL Results
BANDIT Results
Conclusions

INPUT: labeled data

OUTPUT: learned hypothesis used to predict new instances

$h_\theta(x)$, for fixed θ_0 and θ_1 line coefficients

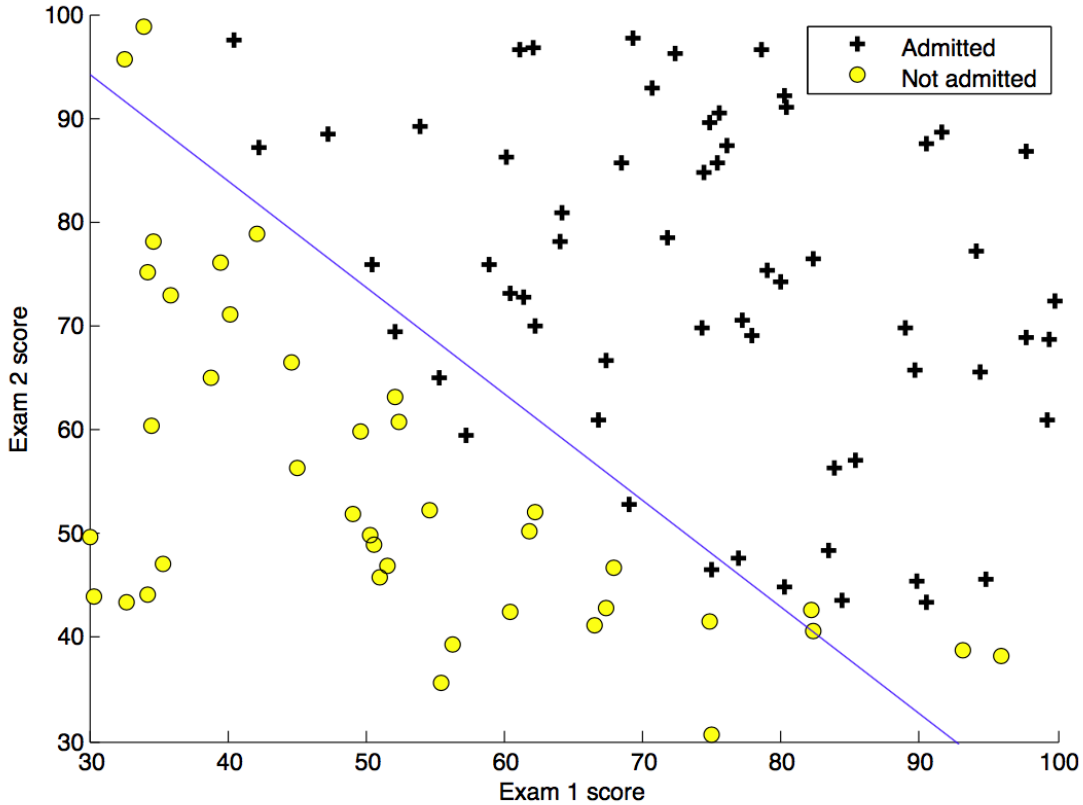

HAL - Protein
James Duin
Introduction
Background
Exp. Setup
Conv. ML
Act. vs Pass.
HAL Results
BANDIT Results
Conclusions

$J(\theta_0, \theta_1)$

(function of the parameters θ_0, θ_1)

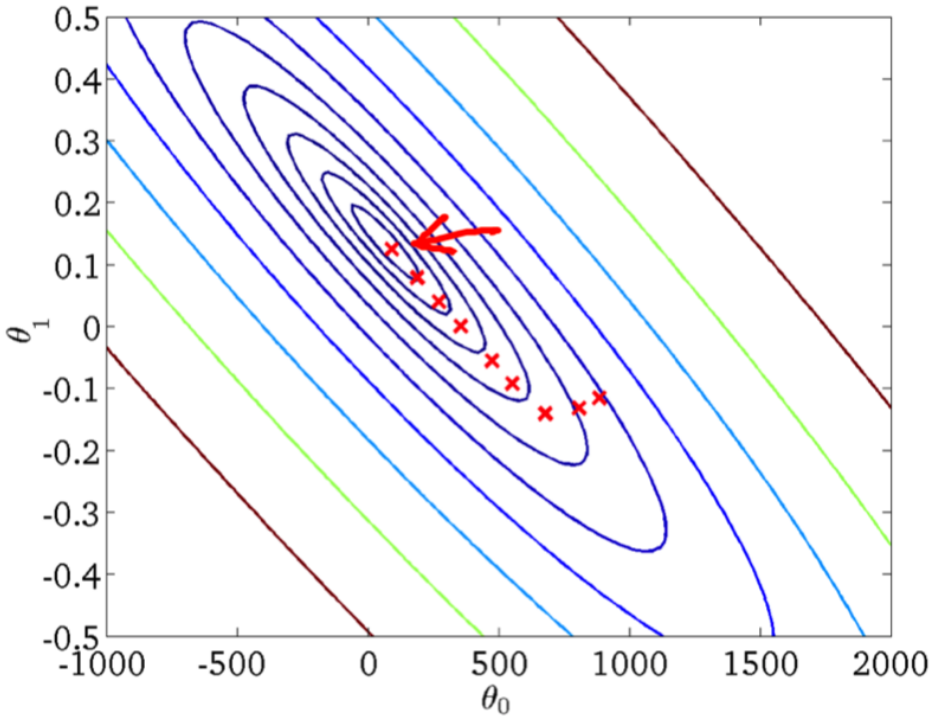

HAL - Protein
James Duin
Introduction
Background
Exp. Setup
Conv. ML
Act. vs Pass.
HAL Results
BANDIT Results
Conclusions

And SVM constructs a line, plane or hyperplane that separates the features with the greatest margin.

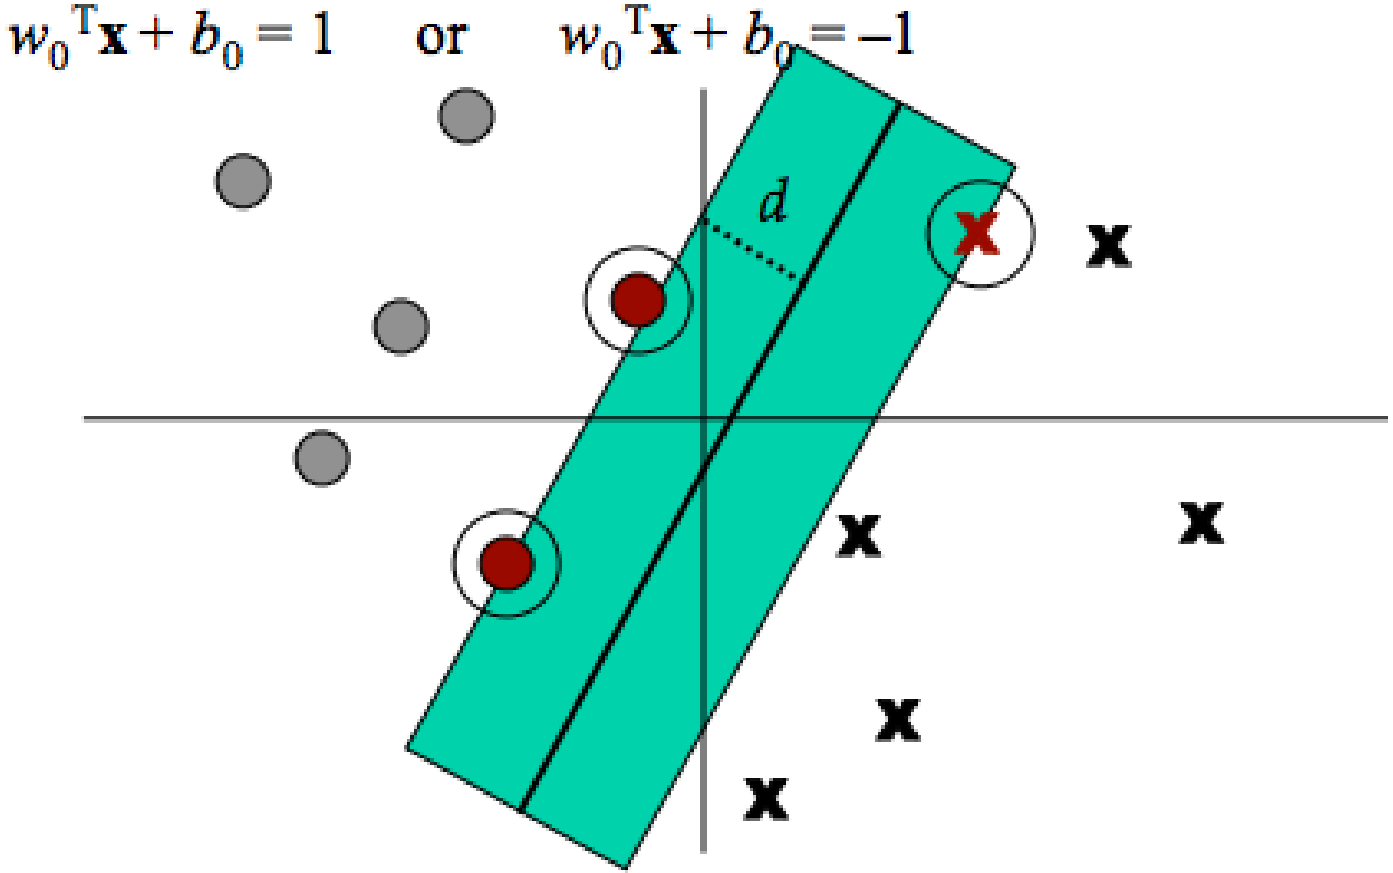

HAL - Protein

James Duin

Introduction

Background

Exp. Setup

Conv. ML

Act. vs Pass.

HAL Results

BANDIT Results

Conclusions

The greater the functional margin the lower the generalization error of the classifier

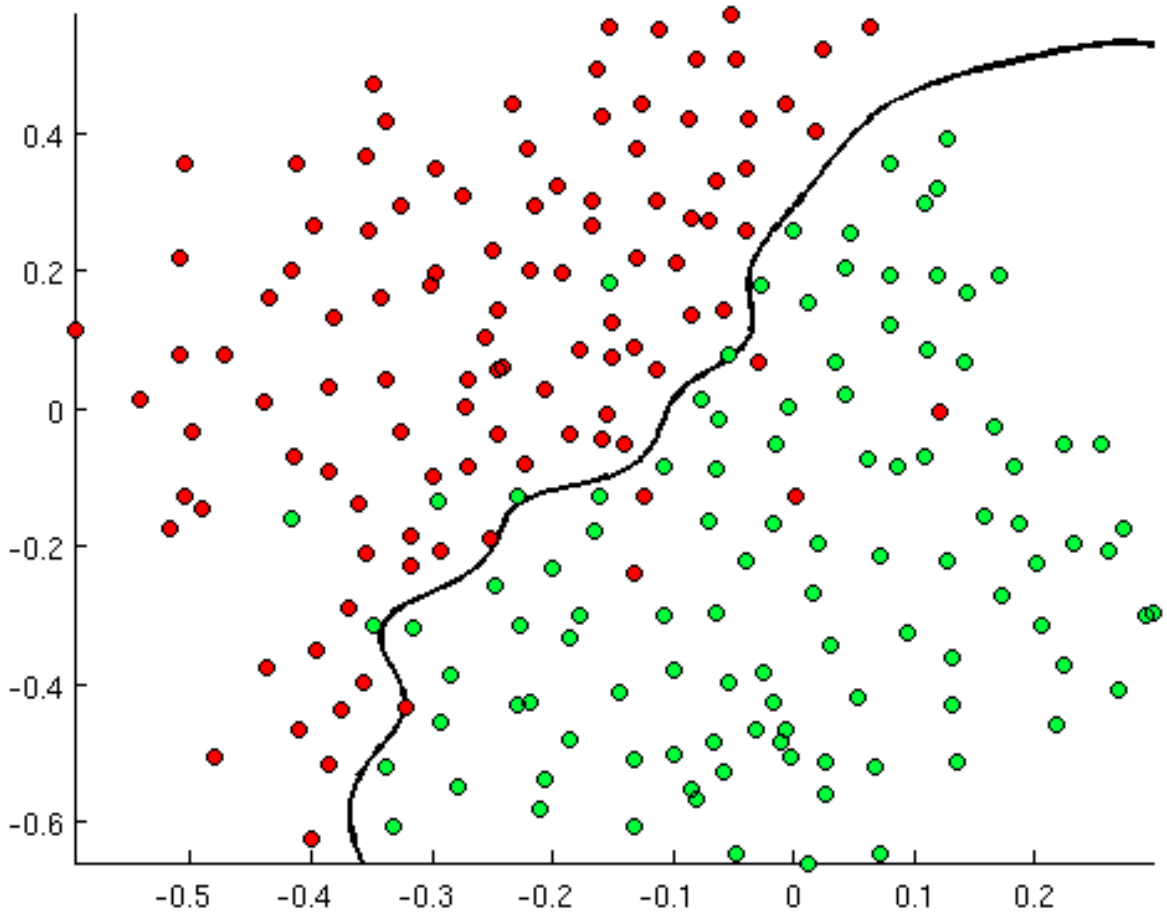

HAL - Protein

James Duin

Introduction

Background

Exp. Setup

Conv. ML

Act. vs Pass.

HAL Results

BANDIT Results

Conclusions

Kernel functions implicitly map inputs into high-dimensional feature spaces

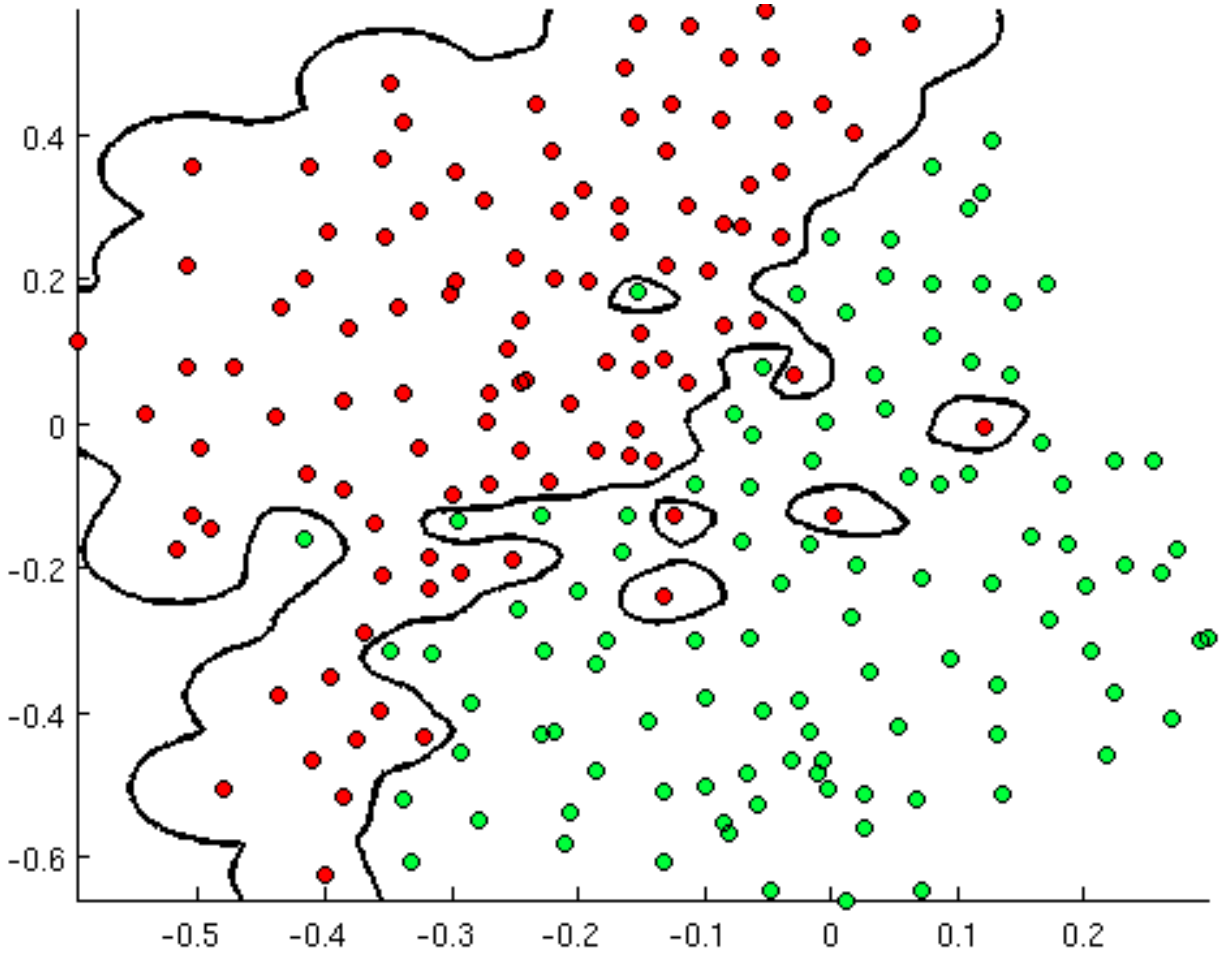

HAL - Protein

James Duin

Introduction

Background

Exp. Setup

Conv. ML

Act. vs Pass.

HAL Results

BANDIT Results

Conclusions

Logistic Regression (Logit) estimates the probability of a binary response, learns coefficients **w** of the input vector **x** and passes dot product through sigmoid function. (Maximum likelihood learning)

$$g(\mathbf{x},\mathbf{w}) = \frac{1}{1+\exp(-\mathbf{w}\cdot\mathbf{x})}$$

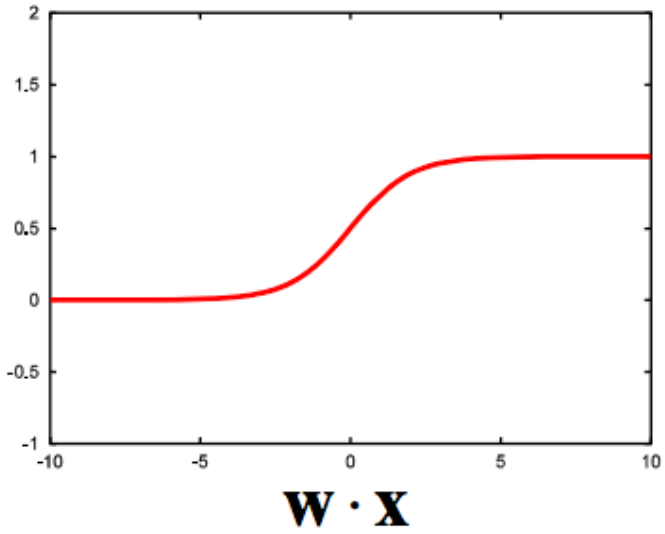

Active Machine Learning

HAL - Protein
James Duin
Introduction
Background
Exp. Setup
Conv. ML
Act. vs Pass.
HAL Results
BANDIT Results
Conclusions

- Unlabeled data is abundant but manually labeling is expensive, e.g. text categorization, drug discovery
- The learner queries an **oracle** or **supervisor** which labels the data at a certain cost
- Active learning solicits new instances that can maximally improve performance of the learned classifier, uncertainty sampling
- Learns the best performing classifier for the minimal amount of labeling cost, or for a given purchase budget
- Acquires labels for each level of the hierachy at a certain cost, spends according to a purchase budget

Active over-labeling
Coarse-grained vs Fine-grained Trade Off

HAL - Protein
James Duin
Introduction
Background
Exp. Setup
Conv. ML
Act. vs Pass.
HAL Results
BANDIT Results
Conclusions

Active over-labeling solicits labels at a finer level of granularity than the target concept

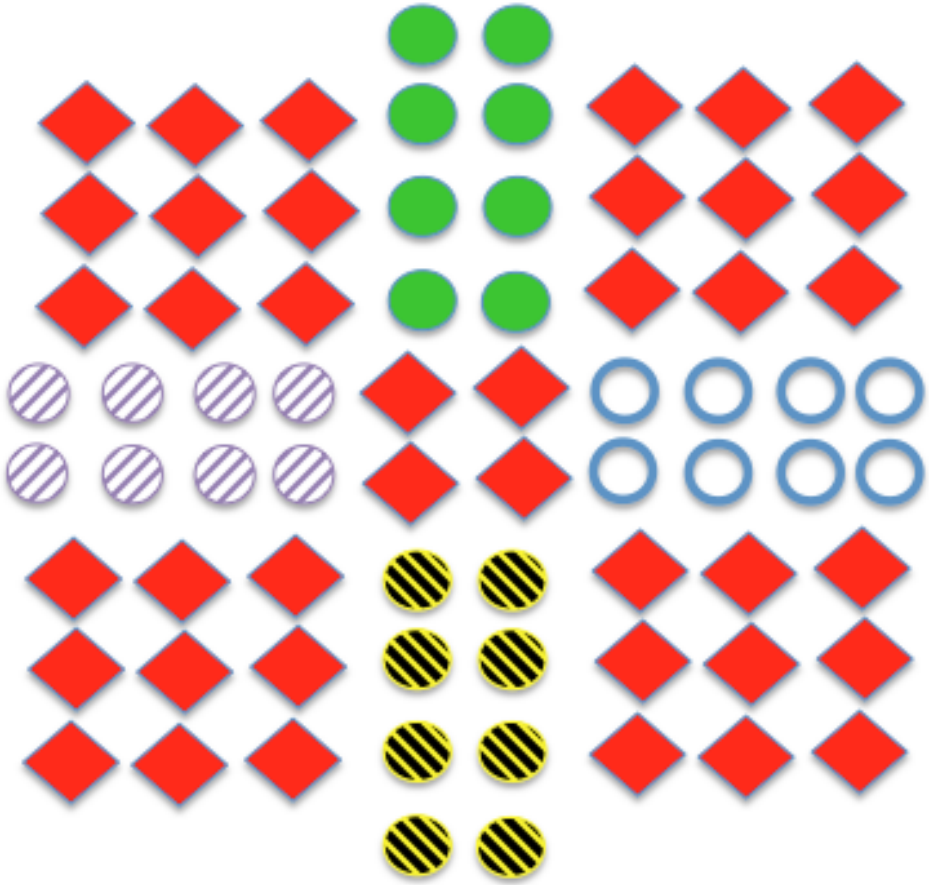

Hierarchical Active Learning

HAL - Protein
James Duin
Introduction
Background
Exp. Setup
Conv. ML
Act. vs Pass.
HAL Results
BANDIT Results
Conclusions

INPUT: purchase proportion p

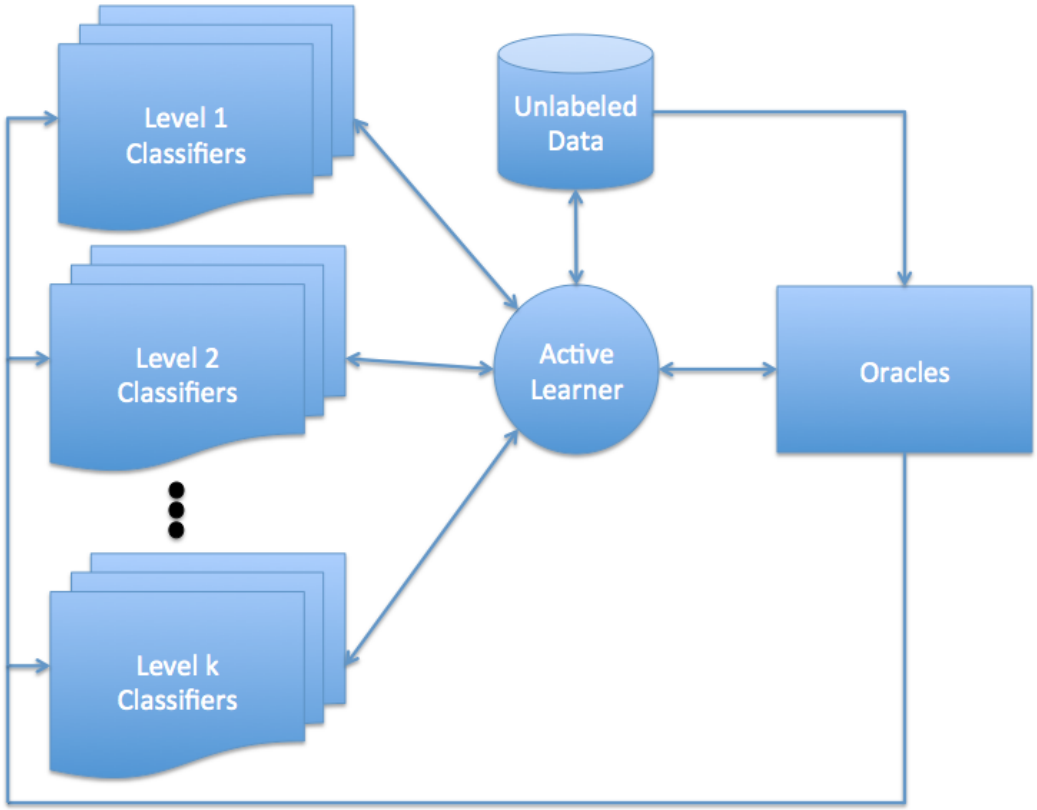

Dynamically Adapting Purchase Proportions

HAL - Protein
James Duin
Introduction
Background
Exp. Setup
Conv. ML
Act. vs Pass.
HAL Results
BANDIT Results
Conclusions

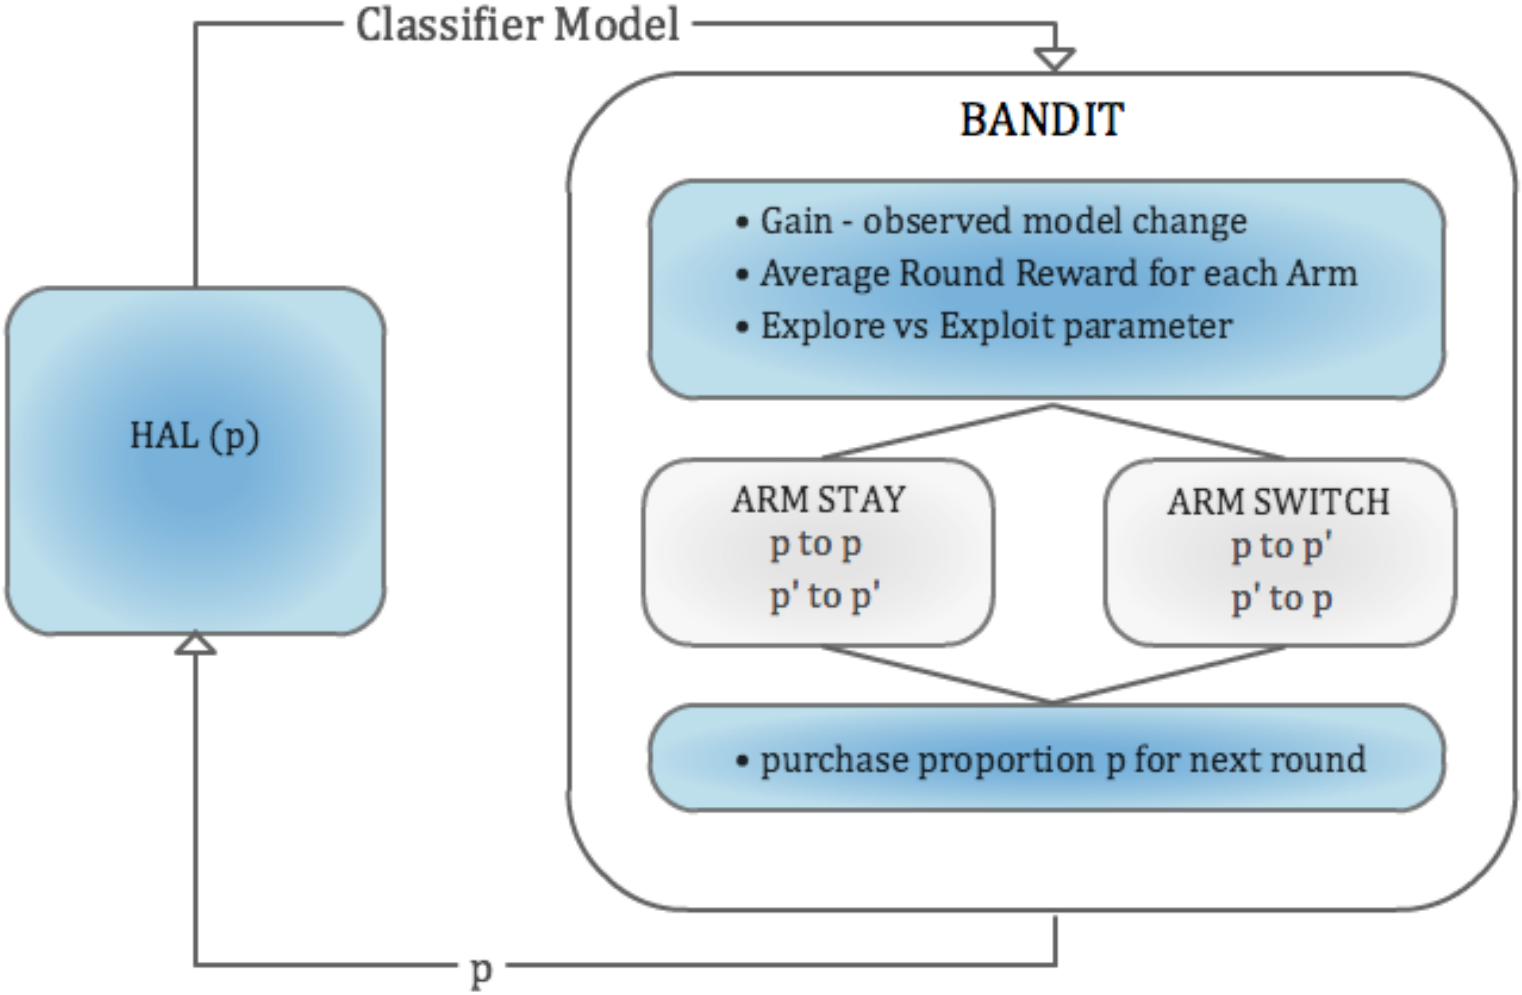

Hierarchical Bioinformatics Data Set

Feature Sources

HAL - Protein
James Duin
Introduction
Background
Exp. Setup
Conv. ML
Act. vs Pass.
HAL Results
BANDIT Results
Conclusions

- Mitoproteome: database of human mitochondrial proteins
- SwissProt: database of experimentally validated human proteins

Type of Properties	Features	Sources
General sequence features	Amino acid composition, sequence length, etc.	Cui et al, PROFEAT
Physico chemical properties	Hydrophobicity, polarity, etc.	Cui et al, PROSO, Phoebus
Structural properties	Secondary structural content, shape, etc.	SSCP
Domains and motifs	Signal peptide, transmembrane domains, etc.	SignalP, TMB-Hunt, NetOgly,TatP

Hierarchical Bioinformatics Data Set

Labeling Hierarchy

HAL - Protein
James Duin
Introduction
Background
Exp. Setup
Conv. ML
Act. vs Pass.
HAL Results
BANDIT Results
Conclusions

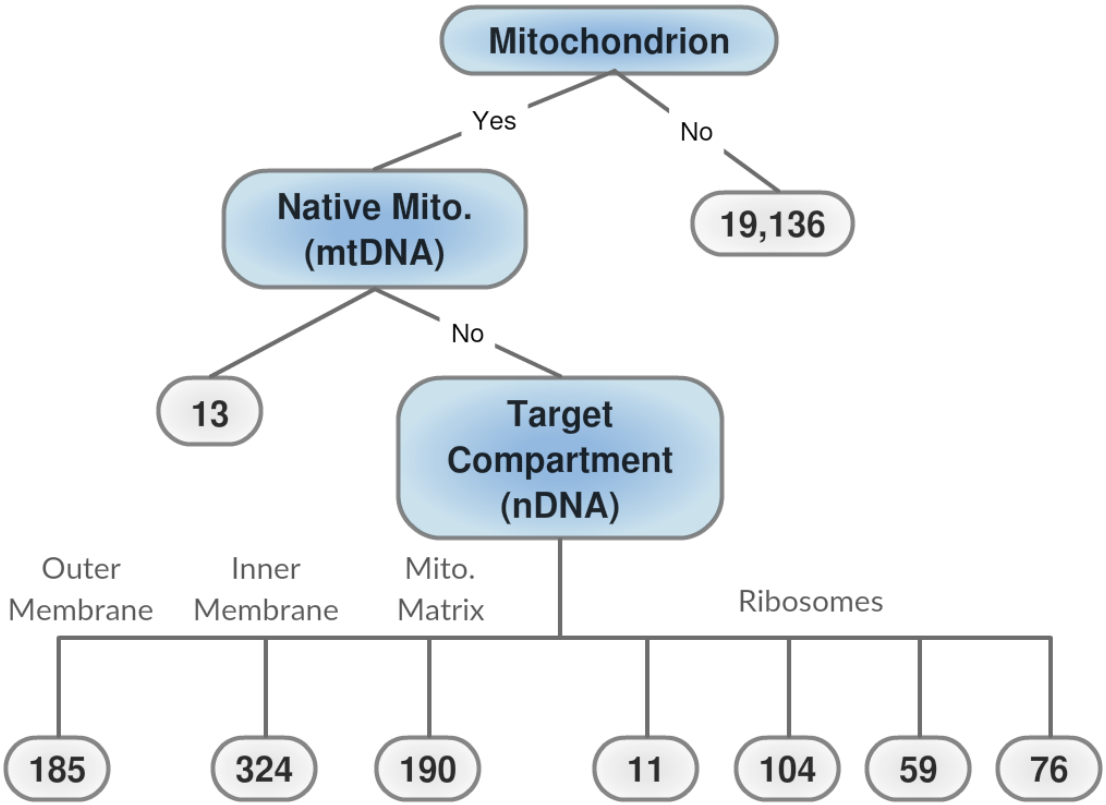

Training and Testing Coarse-Grain and Fine-Grain Classifiers

HAL - Protein
James Duin
Introduction
Background
Exp. Setup
Conv. ML
Act. vs Pass.
HAL Results
BANDIT Results
Conclusions

Number of proteins in each class:

Classes	Count	Totals
Non Mito 0	19136	All: 20098
mtDNA 1	13	Coarse: 19136
nDNA 2	185	Fine: 962
nDNA 3	324	Features: 449
nDNA 4	190	
nDNA 5	11	
nDNA 6	104	
nDNA 7	59	
nDNA 8	76	

SVM and Logit Classifier Performance

Conventional ML

HAL - Protein
James Duin
Introduction
Background
Exp. Setup
Conv. ML
Act. vs Pass.
HAL Results
BANDIT Results
Conclusions

- Tuned parameters for SVM and Logit via an independent run of cross-validation
- **Accuracy**: the percentage of correctly classified results
- **Precision**: a measure of result relevancy
- **Recall**: a measure of how many truly relevant results are returned
- **F-measure**: the harmonic mean of precision and recall
- **PR curve**: plot precision and recall as classifier threshold is varied
- **ROC curve**: plot false positive rate and true positive rate as classifier threshold is varied
- **AUC**: area under the curve, both curves have an optimal AUC of 1.0

SVM and Logit Classifier Performance

F-measure Analysis

HAL - Protein
James Duin
Introduction
Background
Exp. Setup
Conv. ML
Act. vs Pass.
HAL Results
BANDIT Results
Conclusions

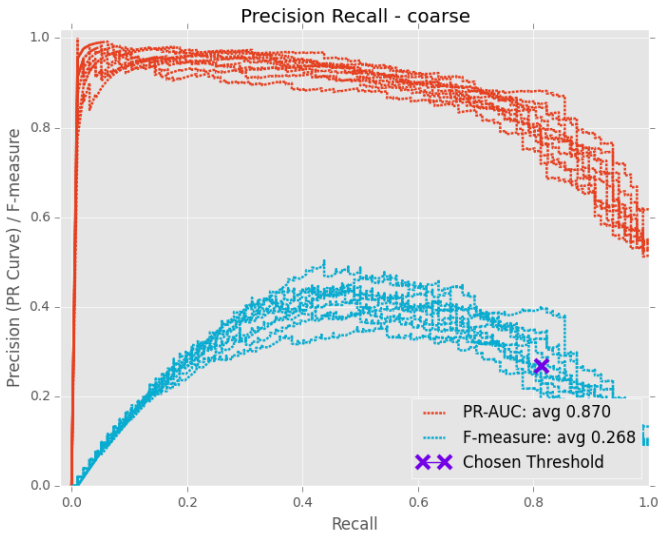

(a) Log Reg Pr Curves - Coarse

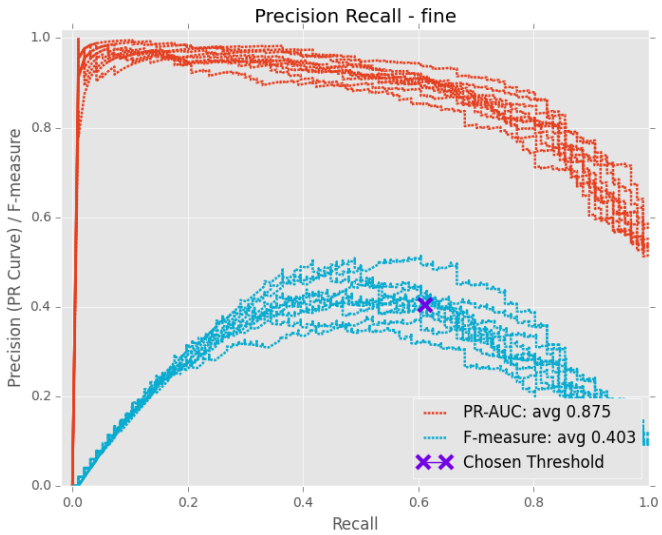

(b) Log Reg Pr Curves - Fine

SVM and Logit Classifier Performance

F-measure Analysis

HAL - Protein
James Duin
Introduction
Background
Exp. Setup
Conv. ML
Act. vs Pass.
HAL Results
BANDIT Results
Conclusions

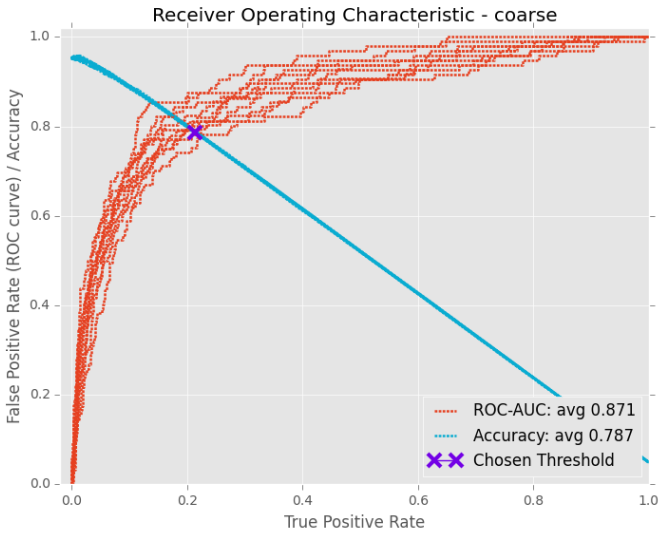

(a) Log Reg ROC Curves - coarse

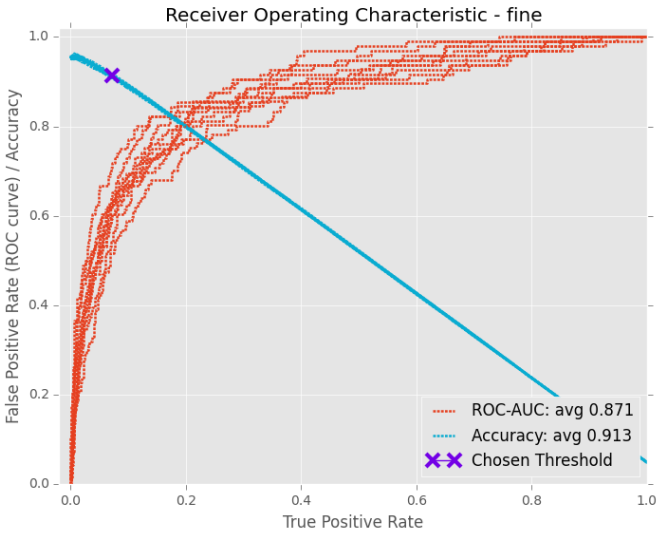

(b) Log Reg ROC Curves - fine

Active vs. Passive Curve Analysis

Logit PR-AUC curves

HAL - Protein
James Duin
Introduction
Background
Exp. Setup
Conv. ML
Act. vs Pass.
HAL Results
BANDIT Results
Conclusions

Iteration: a cycle of the HAL algorithm, a single round of purchasing labels

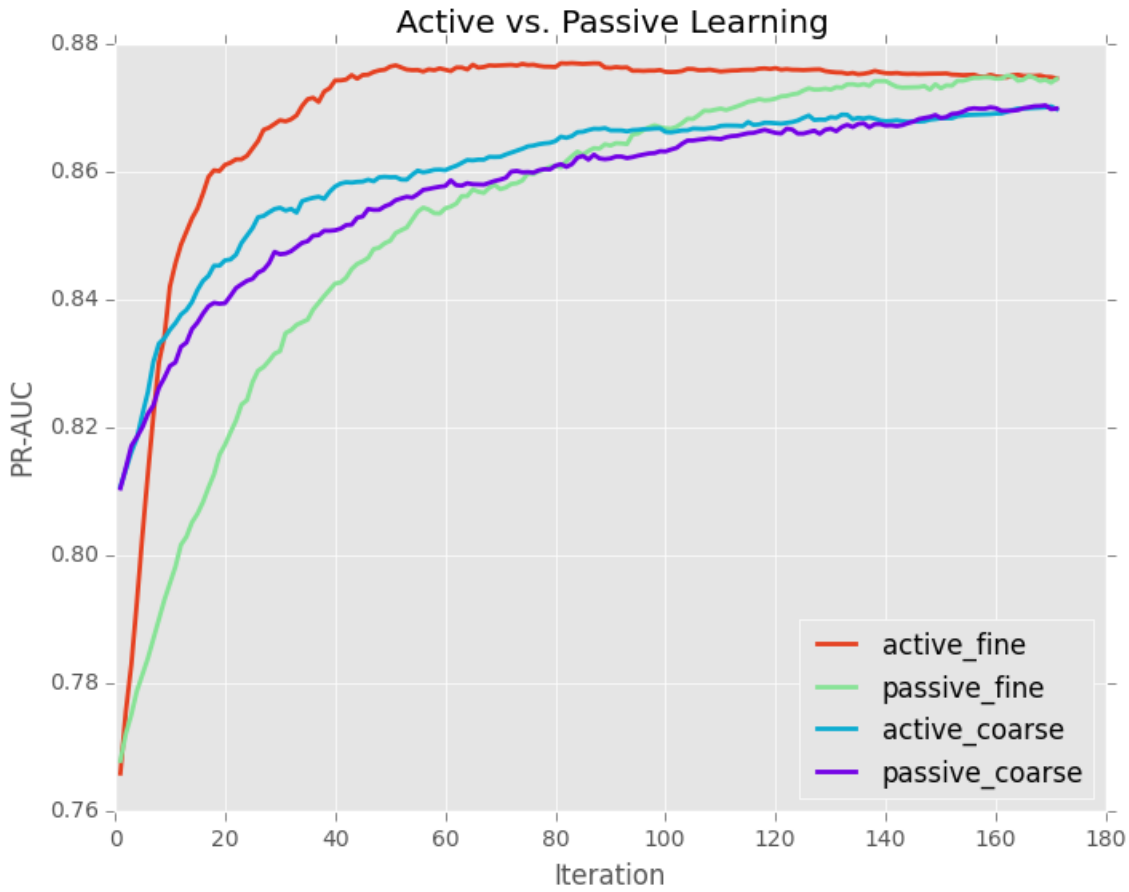

Active vs. Passive Curve Analysis

Logit ROC-AUC curves

HAL - Protein
James Duin
Introduction
Background
Exp. Setup
Conv. ML
Act. vs Pass.
HAL Results
BANDIT Results
Conclusions

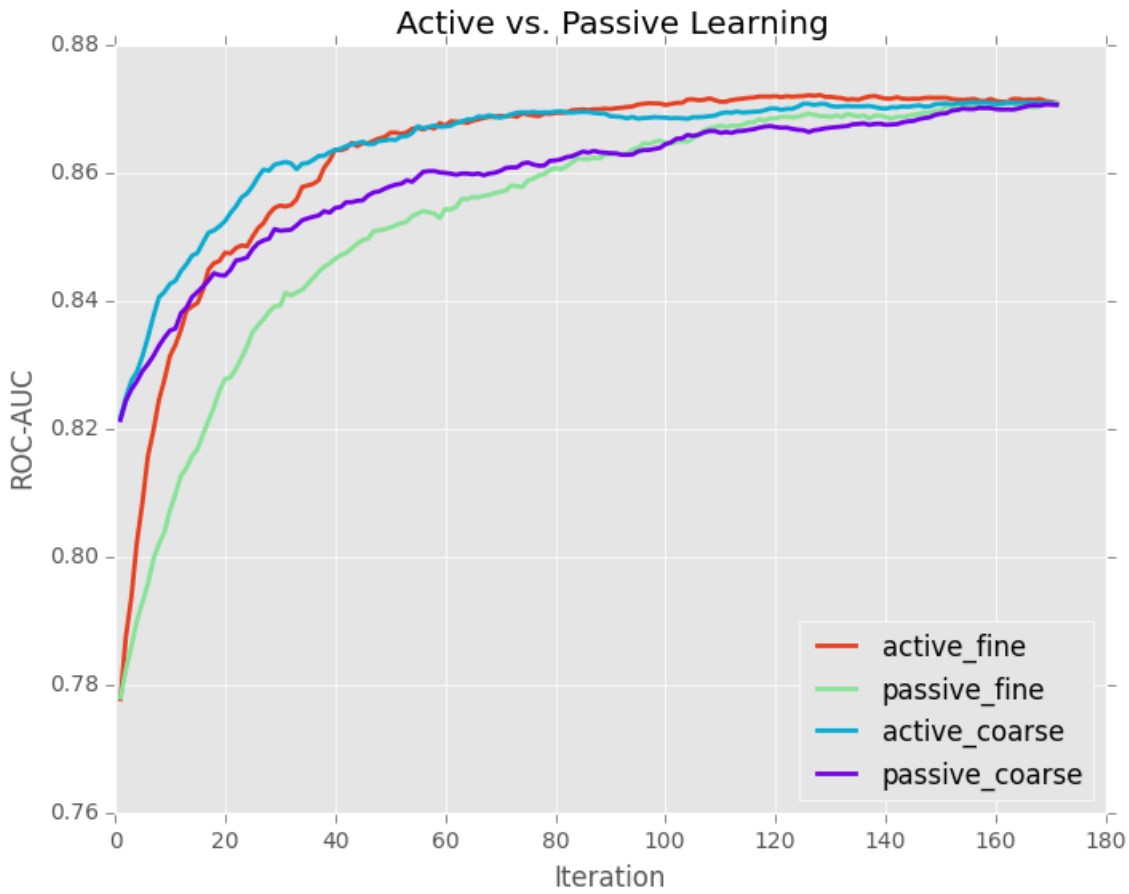

Active vs. Passive Curve Analysis

SVM PR-AUC curves

HAL - Protein
James Duin
Introduction
Background
Exp. Setup
Conv. ML
Act. vs Pass.
HAL Results
BANDIT Results
Conclusions

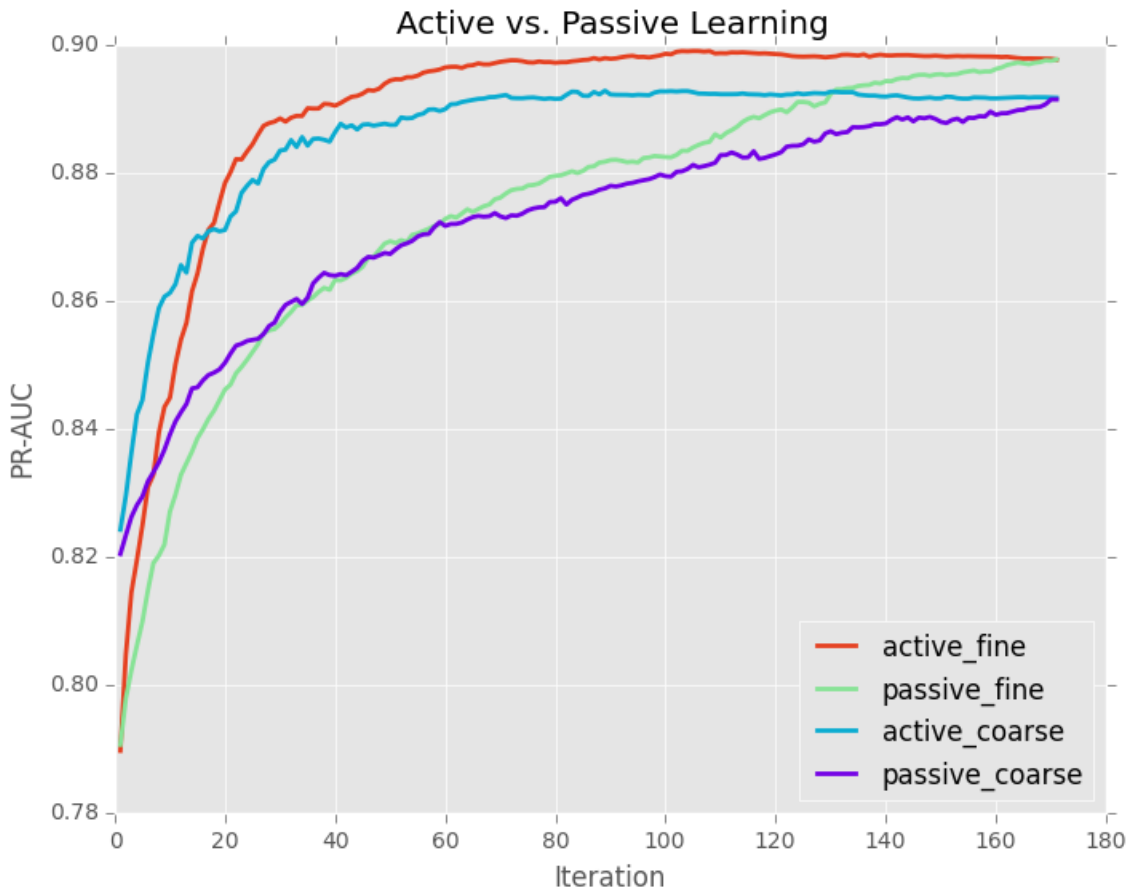

Active vs. Passive Curve Analysis

SVM ROC-AUC curves

HAL - Protein
James Duin
Introduction
Background
Exp. Setup
Conv. ML
Act. vs Pass.
HAL Results
BANDIT Results
Conclusions

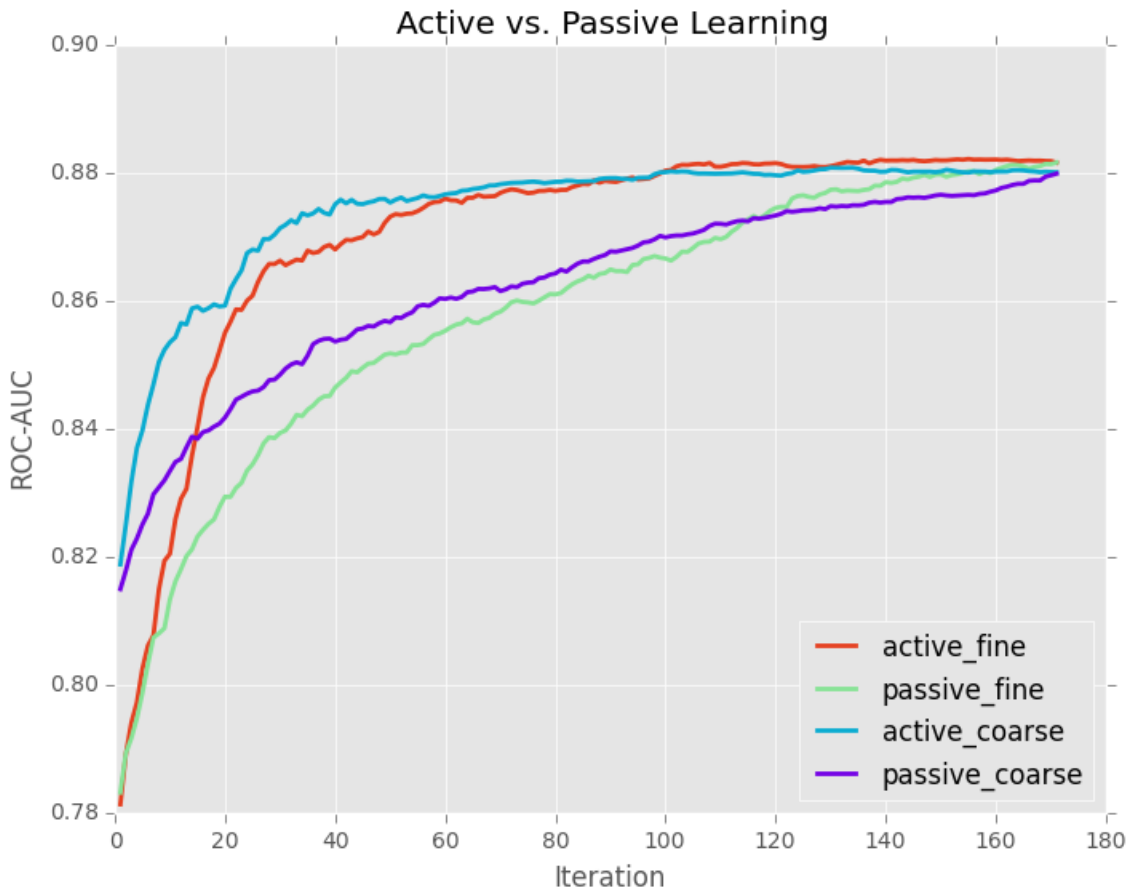

Plots for Fine Fixed Ratio Results

Successive iterations of HAL with fine cost of 1 and coarse cost of 1

HAL - Protein
James Duin
Introduction
Background
Exp. Setup
Conv. ML
Act. vs Pass.
HAL Results
BANDIT Results
Conclusions

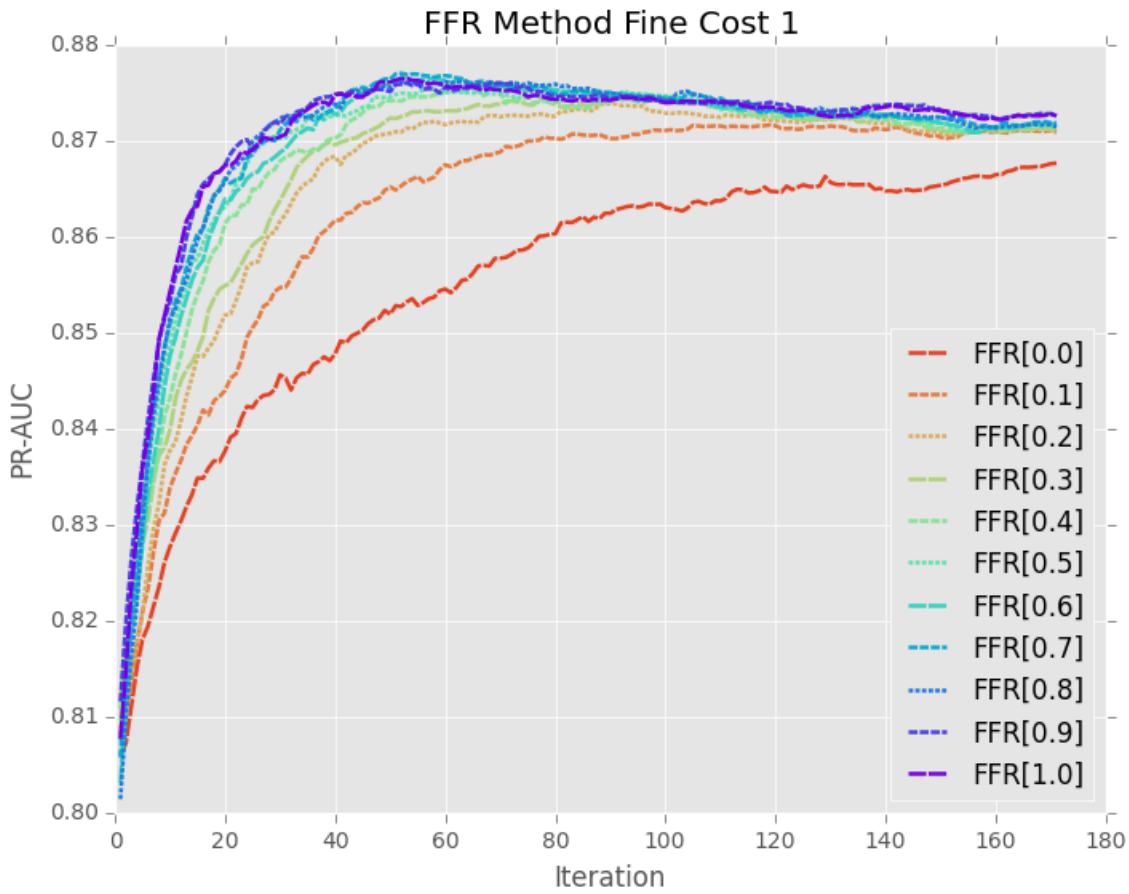

Plots for Fine Fixed Ratio Results

Successive iterations of HAL with fine cost of 4

HAL - Protein
James Duin
Introduction
Background
Exp. Setup
Conv. ML
Act. vs Pass.
HAL Results
BANDIT Results
Conclusions

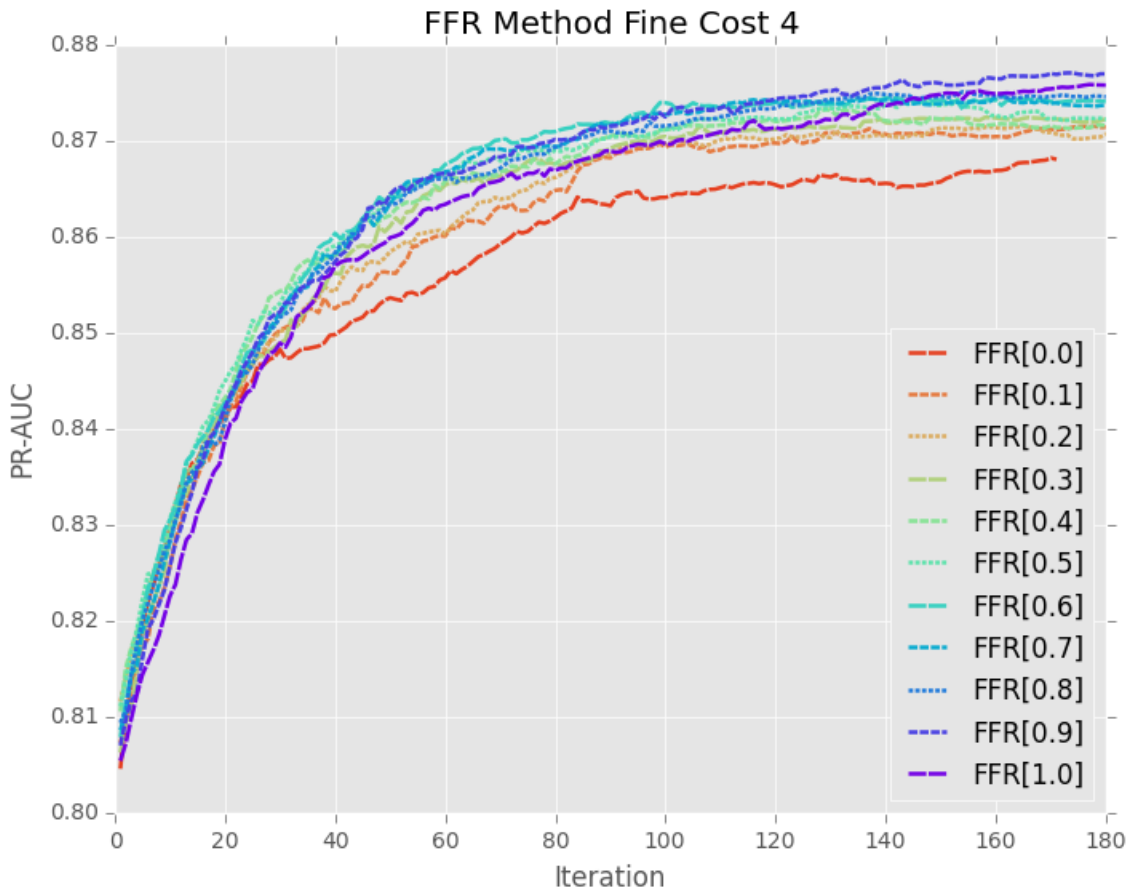

Plots for Fine Fixed Ratio Results

Successive iterations of HAL with fine cost of 8

HAL - Protein
James Duin
Introduction
Background
Exp. Setup
Conv. ML
Act. vs Pass.
HAL Results
BANDIT Results
Conclusions

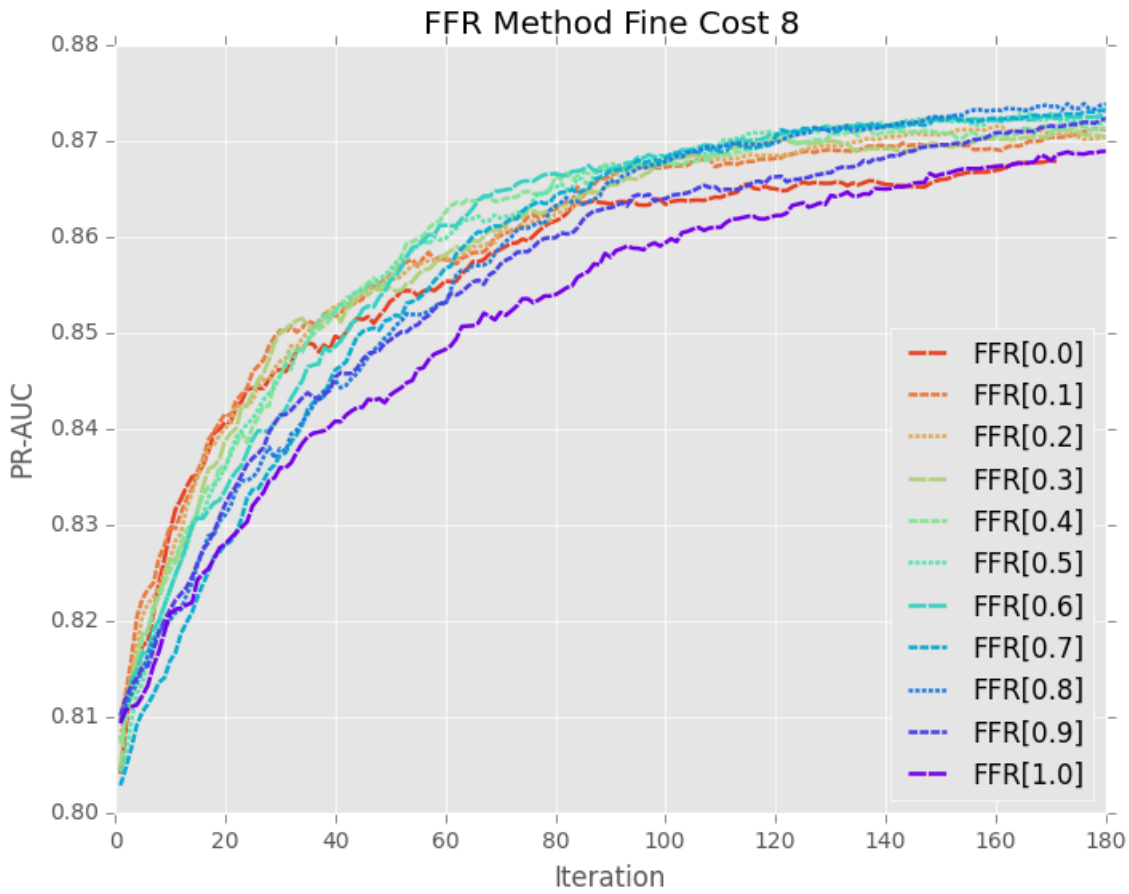

Plots for Fine Fixed Ratio Results

Expanded view Fine Cost 8 - Rnds 20 to 60

HAL - Protein
James Duin
Introduction
Background
Exp. Setup
Conv. ML
Act. vs Pass.
HAL Results
BANDIT Results
Conclusions

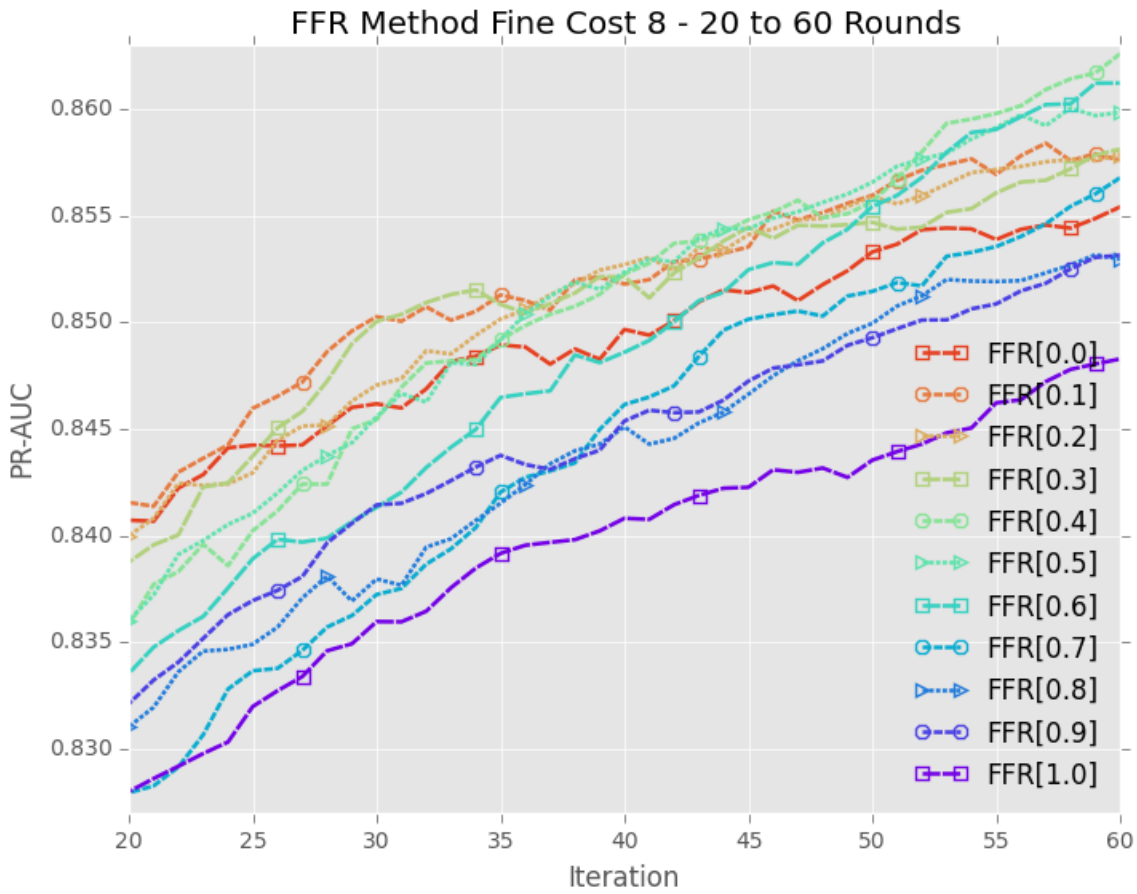

Plots for Fine Fixed Ratio Results

Successive iterations of HAL with fine cost of 16

HAL - Protein
James Duin
Introduction
Background
Exp. Setup
Conv. ML
Act. vs Pass.
HAL Results
BANDIT Results
Conclusions

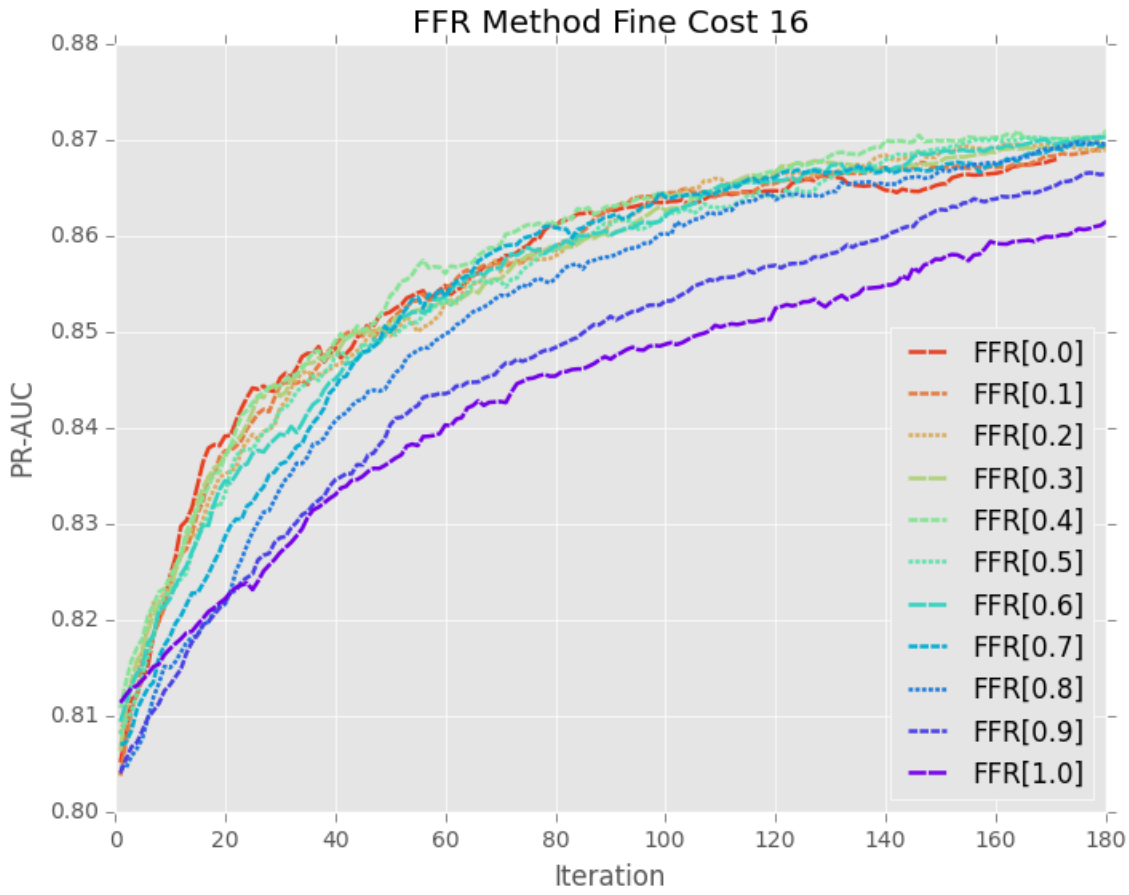

Figure: The fine cost is increased to 16. The fine cost is to high to offset the decreased number of instances purchased.

BANDIT Approach Results
Varying Cost Analysis

HAL - Protein
James Duin
Introduction
Background
Exp. Setup
Conv. ML
Act. vs Pass.
HAL Results
BANDIT Results
Conclusions

- The BANDIT approach is compared to the previous FFR curves for the following fine-grain costs $\{1.0, 1.1, 1.2, 1.5, 2.0, 4.0, 8.0, 16.0, 32.0, 64.0\}$
- Budget held fixed at round 120.
- The metric *diff* is the learner's absolute difference in PR-AUC from the top learner for a given cost.
- The metric *rank* is the learners 0 indexed ranking in terms of PR-AUC for a given cost.

BANDIT Approach Results

Varying Cost Analysis - Plot

HAL - Protein
James Duin
Introduction
Background
Exp. Setup
Conv. ML
Act. vs Pass.
HAL Results
BANDIT Results
Conclusions

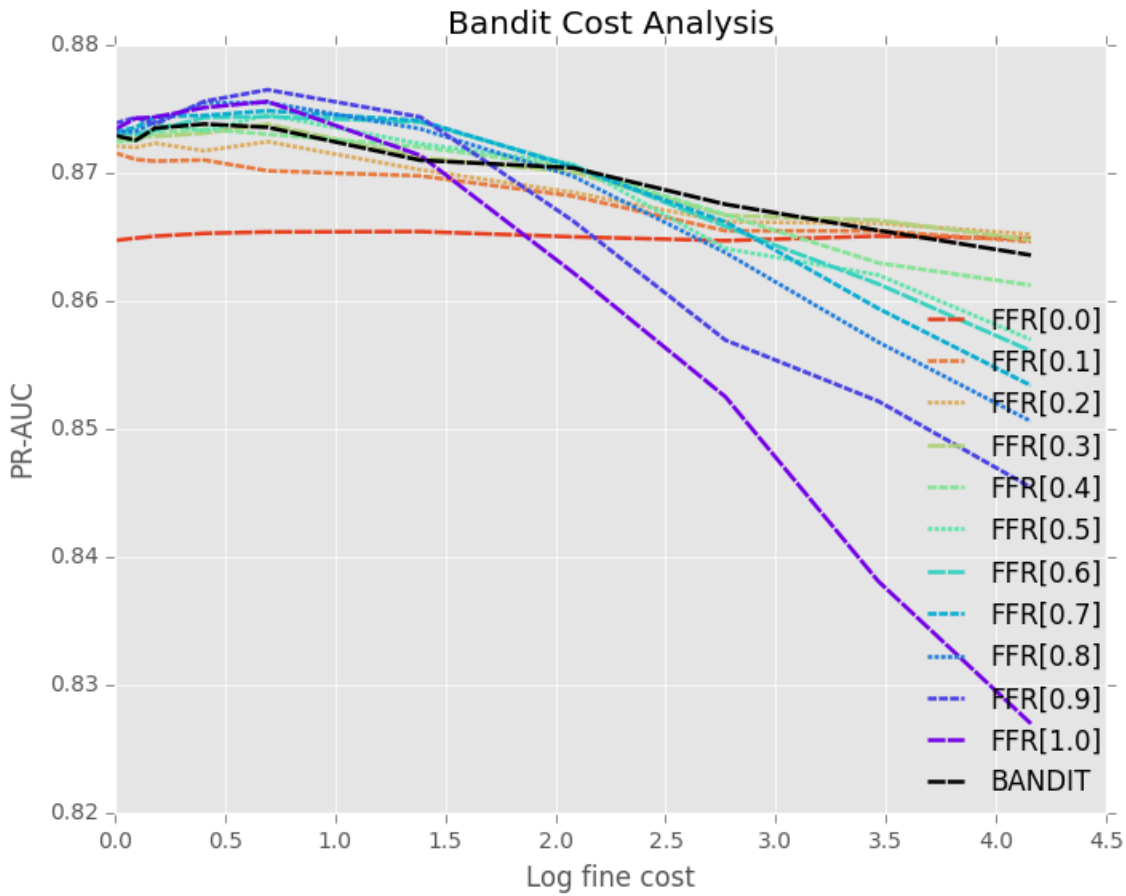

BANDIT Approach Results

Varying Cost Analysis - Rank and Diff Metrics

HAL - Protein
James Duin
Introduction
Background
Exp. Setup
Conv. ML
Act. vs Pass.
HAL Results
BANDIT Results
Conclusions

	diff				rank			
	min	max	mean	std	min	max	mean	std
algorithm								
BANDIT	0.000	0.003	**0.001**	0.001	0	8	4.8	2.315
FFR[0.0]	0.000	0.011	0.007	0.004	1	11	8.8	3.429
FFR[0.1]	0.001	0.006	0.003	0.002	3	10	8.0	2.793
FFR[0.2]	0.000	0.004	0.002	0.001	0	9	6.5	3.500
FFR[0.3]	0.000	0.003	0.001	0.001	0	8	5.1	2.663
FFR[0.4]	0.000	0.004	0.002	0.001	1	8	5.6	2.200
FFR[0.5]	0.000	0.008	0.002	0.002	0	8	4.6	2.200
FFR[0.6]	0.000	0.009	0.002	0.003	1	7	4.6	1.855
FFR[0.7]	0.000	0.012	0.002	0.004	0	8	**3.3**	2.571
FFR[0.8]	0.000	0.015	0.003	0.005	1	9	4.8	3.027
FFR[0.9]	0.000	0.020	0.005	0.007	0	10	4.3	4.605
FFR[1.0]	0.000	0.038	0.009	0.013	1	11	5.6	4.630

BANDIT Approach Results

Varying Budget Analysis - Mixed Cost

HAL - Protein
James Duin
Introduction
Background
Exp. Setup
Conv. ML
Act. vs Pass.
HAL Results
BANDIT Results
Conclusions

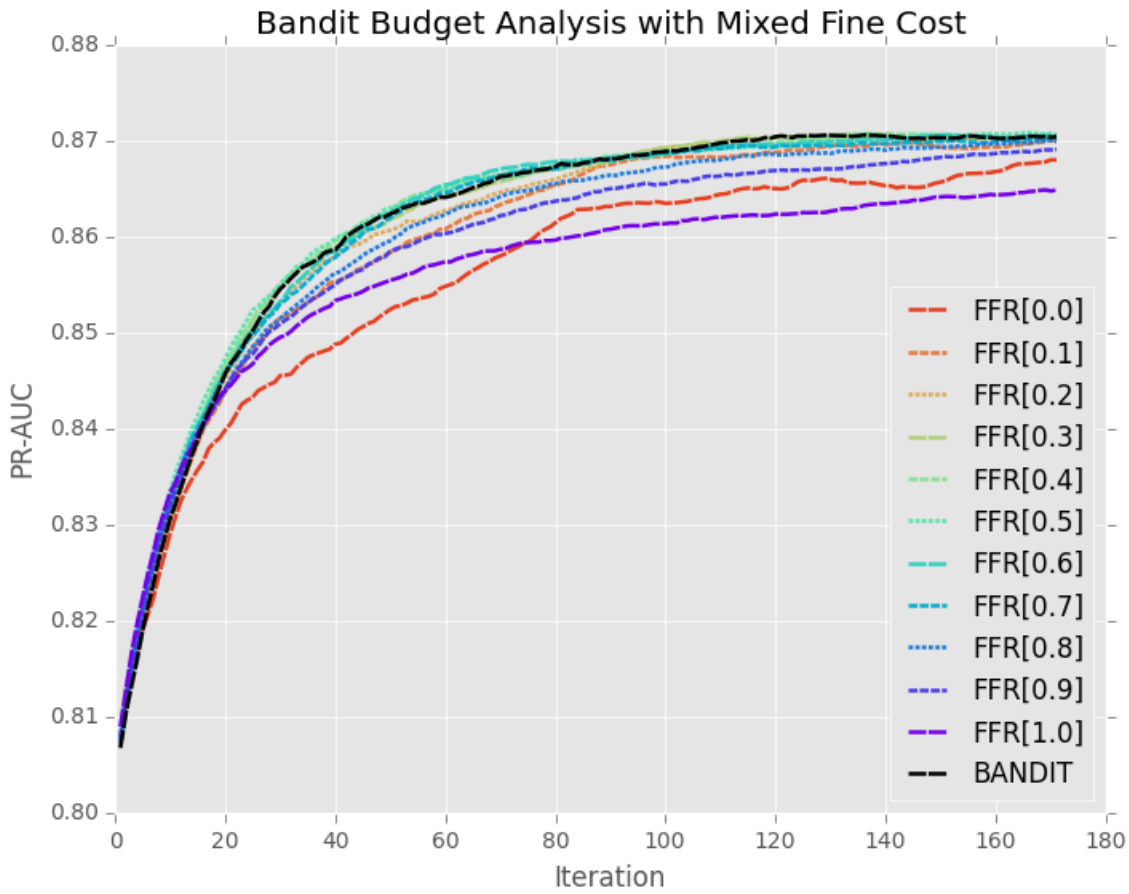

BANDIT Approach Results

BANDIT - Rnds 20 to 60

HAL - Protein
James Duin
Introduction
Background
Exp. Setup
Conv. ML
Act. vs Pass.
HAL Results
BANDIT Results
Conclusions

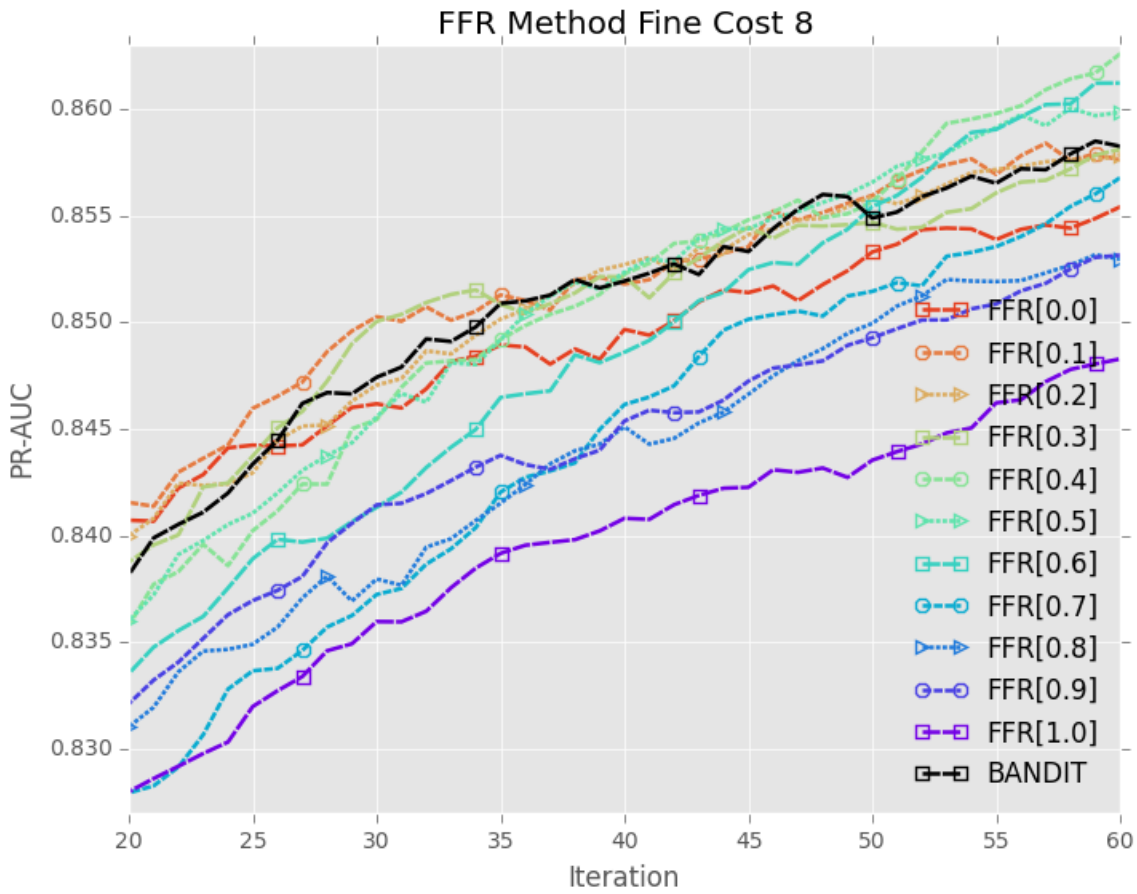

Conclusions

- Demonstrated fine-grained labels can be used to improve a coarse-grained classifier for the protein dataset
- Demonstrated a prominent advantage for active fine with the Logit classifier
- HAL is implemented and applied to the protein dataset for various FFR proportions and fine label costs
- The BANDIT approach is shown to be robust to both labeling cost and budget

Future Work

HAL - Protein
James Duin
Introduction
Background
Exp. Setup
Conv. ML
Act. vs Pass.
HAL Results
BANDIT Results
Conclusions

- Future work is to apply the active over-labeling approach to other datasets with more complex hierarchical label trees; datasets derived from Gene Ontology research could be investigated

Acknowledgements

HAL - Protein
James Duin
Introduction
Background
Exp. Setup
Conv. ML
Act. vs Pass.
HAL Results
BANDIT Results
Conclusions

- I would like to thank my advisor Dr. Stephen Scott, Yugi Mo and Dr. Douglas Downey for continued guidance. I would like to thank Dr. Juan Cui and Dr. Ashok Samal for serving on my comittee. Additionally, I would like to thank Jiang Shu and Kevin Chiang for their assistance accessing and understanding the protein dataset.

Questions

jamesdduin@gmail.com

Plots for Fine Fixed Ratio Results

Fine Cost 8 - Rnds to 500

HAL - Protein
James Duin
Introduction
Background
Exp. Setup
Conv. ML
Act. vs Pass.
HAL Results
BANDIT Results
Conclusions

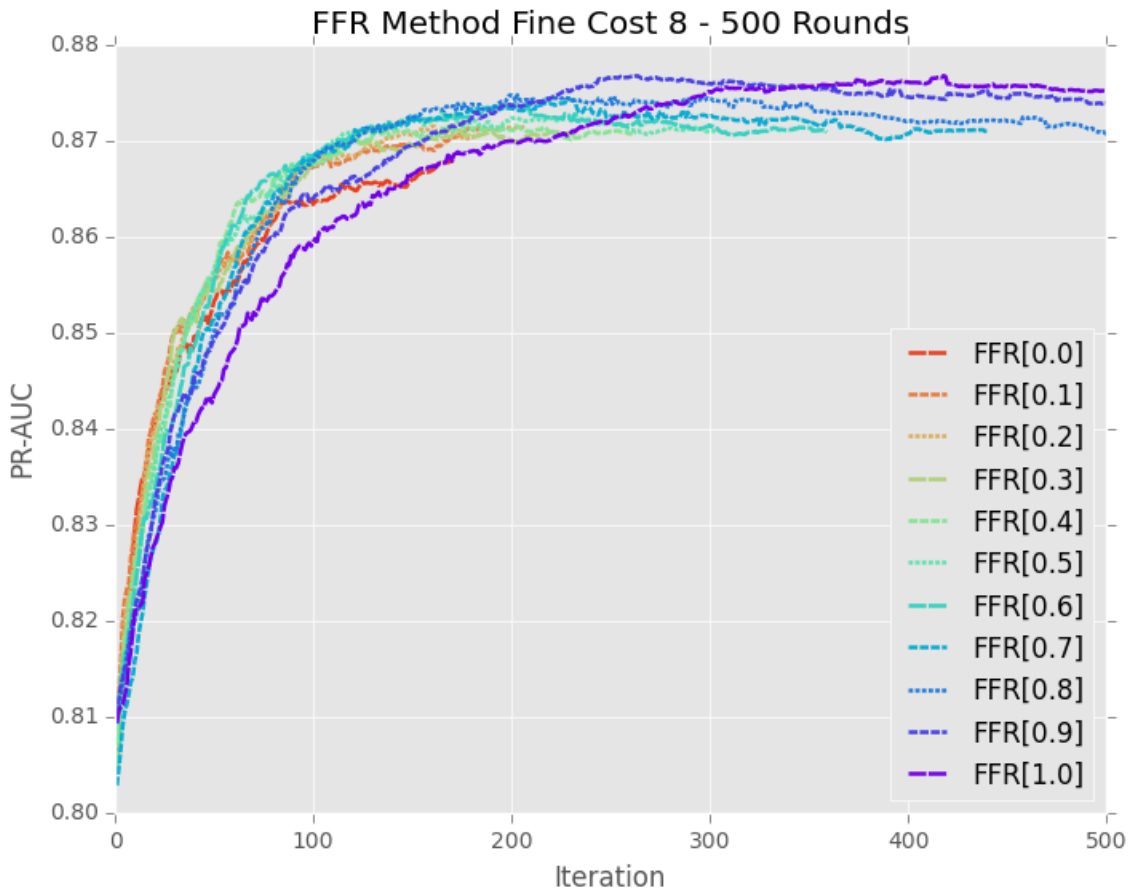

Plots for Fine Fixed Ratio Results

Successive iterations of HAL with fine cost of 2

HAL - Protein
James Duin
Introduction
Background
Exp. Setup
Conv. ML
Act. vs Pass.
HAL Results
BANDIT Results
Conclusions

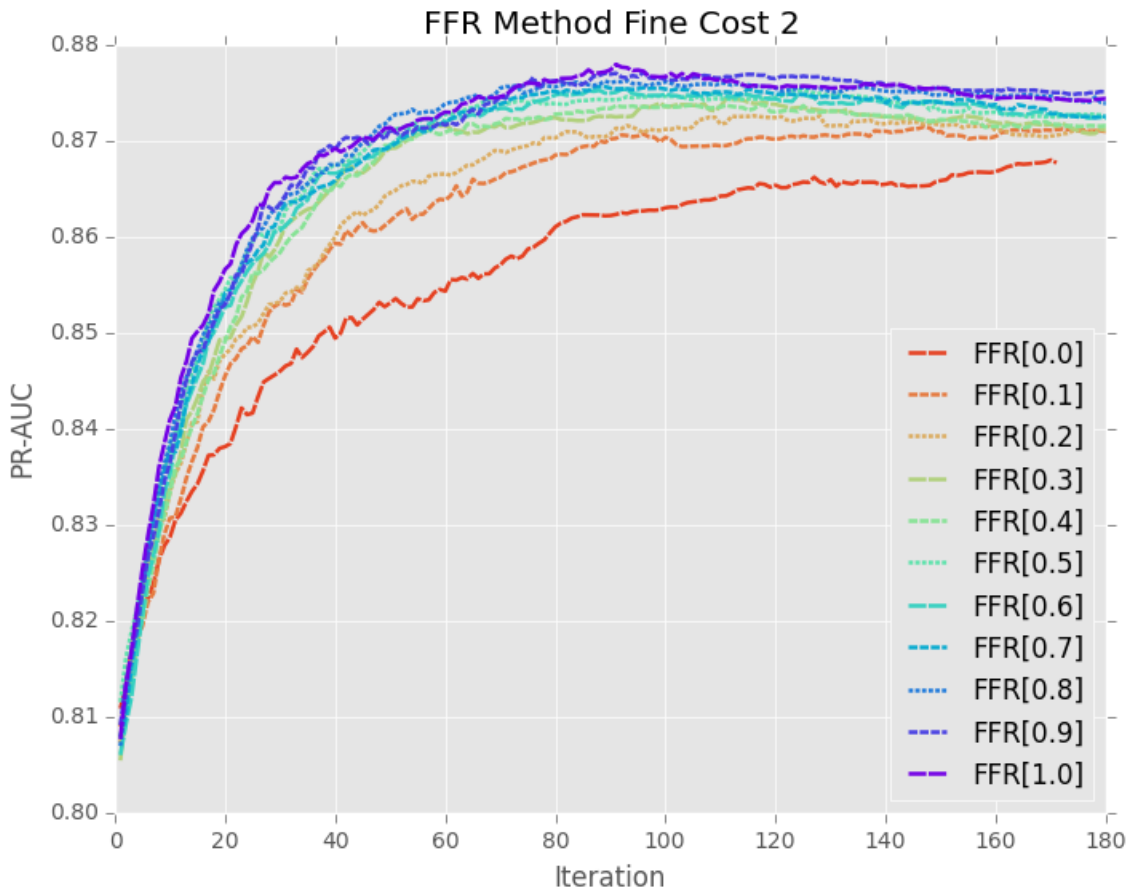

Evaluating Classifier Performance
Confusion Matrix

HAL - Protein
James Duin
Introduction
Background
Exp. Setup
Conv. ML
Act. vs Pass.
HAL Results
BANDIT Results
Conclusions

Divide data into train and a test set. Analyze test set with the following values:

- True-Negatives (T_n): Correctly classified negatives
- False-Negatives (F_p): Incorrectly classified negatives
- False-Positives (F_n): Incorrectly classified positives
- True-Positives (T_p): Correctly classified positives

Example of a confusion matrix for a test set with 100 negatives and 50 positives:

conf (T_n/F_n)	conf (F_p/T_p)
90	10
20	30

Evaluating Classifier Performance
Precision and Recall

HAL - Protein
James Duin
Introduction
Background
Exp. Setup
Conv. ML
Act. vs Pass.
HAL Results
BANDIT Results
Conclusions

Precision is a measure of result relevancy:

$$P = \frac{T_p}{T_p + F_p} \tag{1}$$

Recall is a measure of how many truly relevant results are returned:

$$R = \frac{T_p}{T_p + F_n} \tag{2}$$

HAL - Protein

James Duin

Introduction

Background

Exp. Setup

Conv. ML

Act. vs Pass.

HAL Results

BANDIT Results

Conclusions

The F-measure or F1-measure (F1) is the harmonic mean of precision and recall:

$$F1 = 2 \cdot \frac{P \cdot R}{P + R} \tag{3}$$

Evaluating Classifier Performance
ROC - PR curves

HAL - Protein
James Duin
Introduction
Background
Exp. Setup
Conv. ML
Act. vs Pass.
HAL Results
BANDIT Results
Conclusions

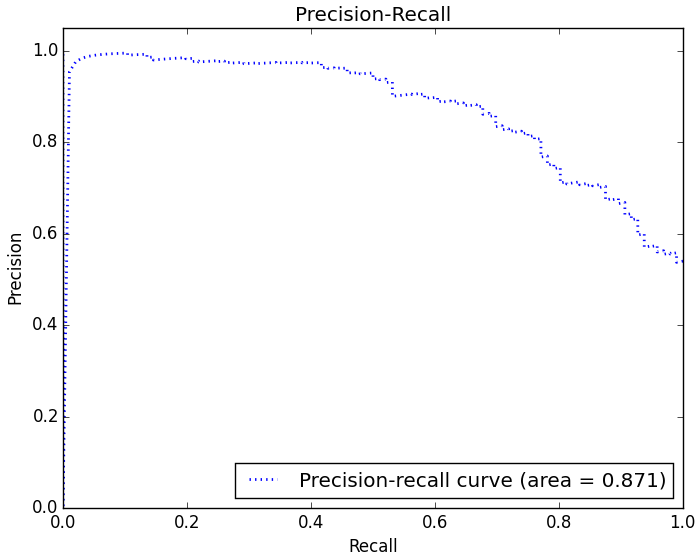

(a) PR curve.

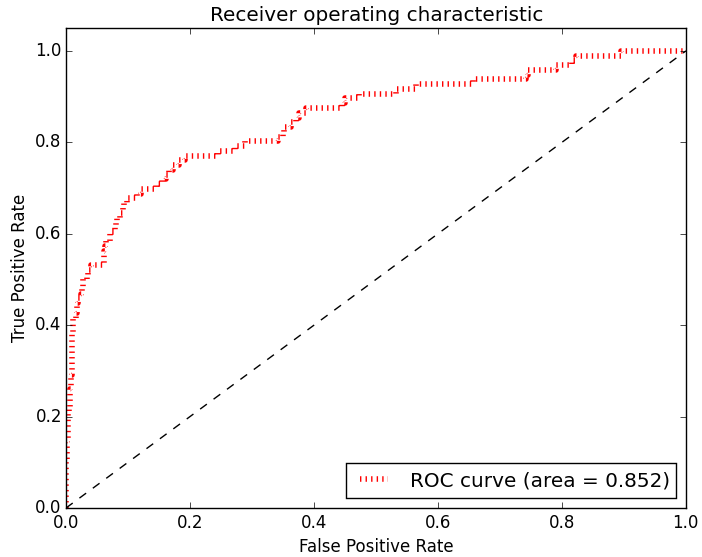

(b) ROC curve.

Figure: Examples of PR and ROC curves with their corresponding AUC values.

Training and Testing Coarse-Grain and Fine-Grain Classifiers

HAL - Protein
James Duin
Introduction
Background
Exp. Setup
Conv. ML
Act. vs Pass.
HAL Results
BANDIT Results
Conclusions

Table: Number of proteins in each partition:

Folds	All	0	1	2	3	4	5	6	7	8
1	2010	1914	1	19	32	19	1	11	6	7
2	2010	1914	1	19	32	19	1	11	6	7
3	2010	1914	1	19	32	19	1	11	5	8
4	2010	1914	1	19	32	19	1	10	6	8
5	2010	1914	1	18	33	19	1	10	6	8
6	2010	1914	1	18	33	19	1	10	6	8
7	2010	1913	2	18	33	19	1	10	6	8
8	2010	1913	2	18	33	19	1	10	6	8
9	2009	1913	2	18	32	19	2	10	6	7
10	2009	1913	1	19	32	19	1	11	6	7
Total	20098	19136	13	185	324	190	11	104	59	76

Table: Logit results after parameter tuning:

Title	PR	ROC	Acc	F1	conf (tn/fn)	conf (fp/tp)
coarse	0.870	0.871	0.787	0.268	(1503.2 / 17.8)	(410.4 / 78.3)
fine	0.875	0.871	0.913	0.403	(1776.5 / 37.3)	(137.1 / 58.8)

Table: SVM results after parameter tuning:

Title	PR	ROC	Acc	F1	conf (tn/fn)	conf (fp/tp)
coarse	0.892	0.880	0.866	0.347	(1669.5 / 24.8)	(244.1 / 71.3)
fine	0.898	0.882	0.942	0.485	(1839.0 / 41.5)	(74.6 / 54.6)

Application to Dispatch Dataset

HAL - Protein
James Duin
Introduction
Background
Exp. Setup
Conv. ML
Act. vs Pass.
HAL Results
BANDIT Results
Conclusions

Analysis and evaluation follow Mo et al.'s work.

- Fine outperforms Coarse in PR-AUC
- Active outperforms Passive in PR-AUC
- HAL ran with variable cost, fine proportions and budget
- BANDIT approach shown to be robust to changes in cost and budget

Active vs. Passive Curve Analysis

Logit Accuracy

HAL - Protein
James Duin
Introduction
Background
Exp. Setup
Conv. ML
Act. vs Pass.
HAL Results
BANDIT Results
Conclusions

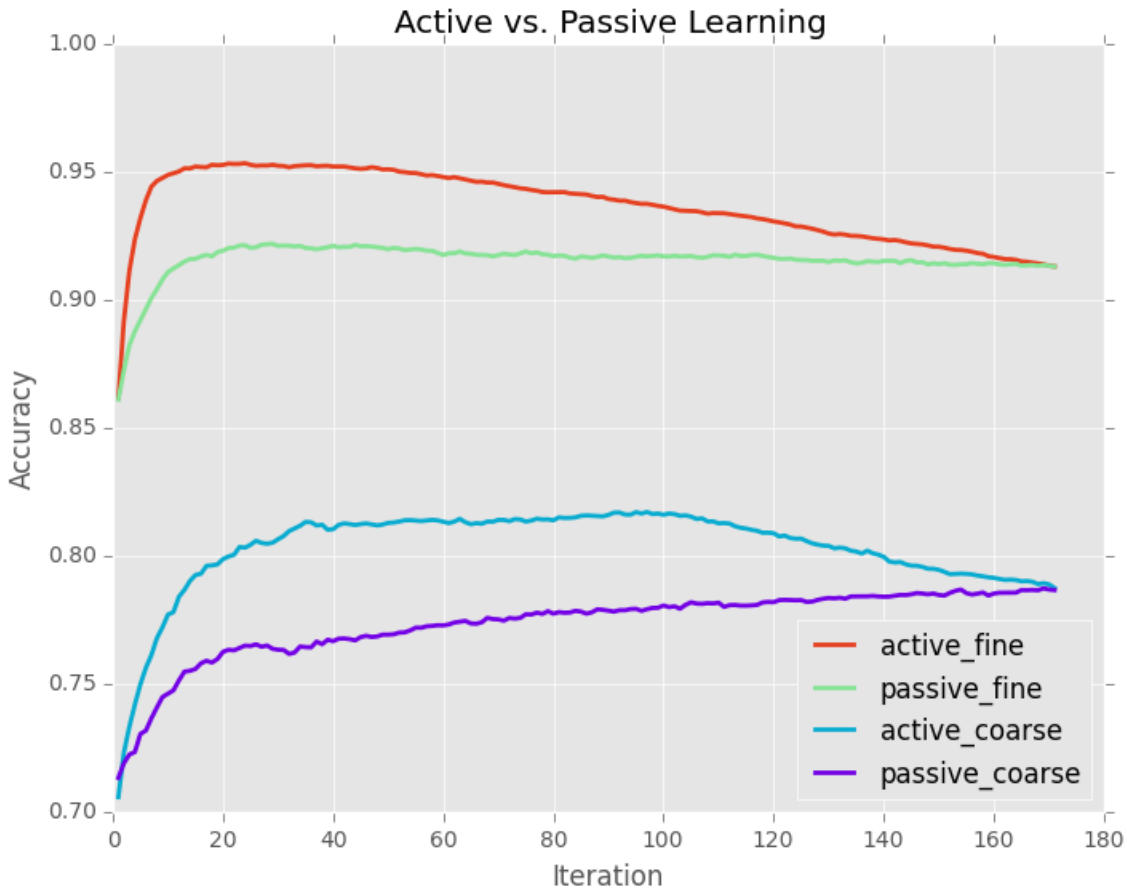

Figure: The accuracy of the classifiers stays at roughly the same rate throughout the rounds; this is due to an effective weighting scheme.

Active vs. Passive Curve Analysis

Logit F-measure

HAL - Protein
James Duin
Introduction
Background
Exp. Setup
Conv. ML
Act. vs Pass.
HAL Results
BANDIT Results
Conclusions

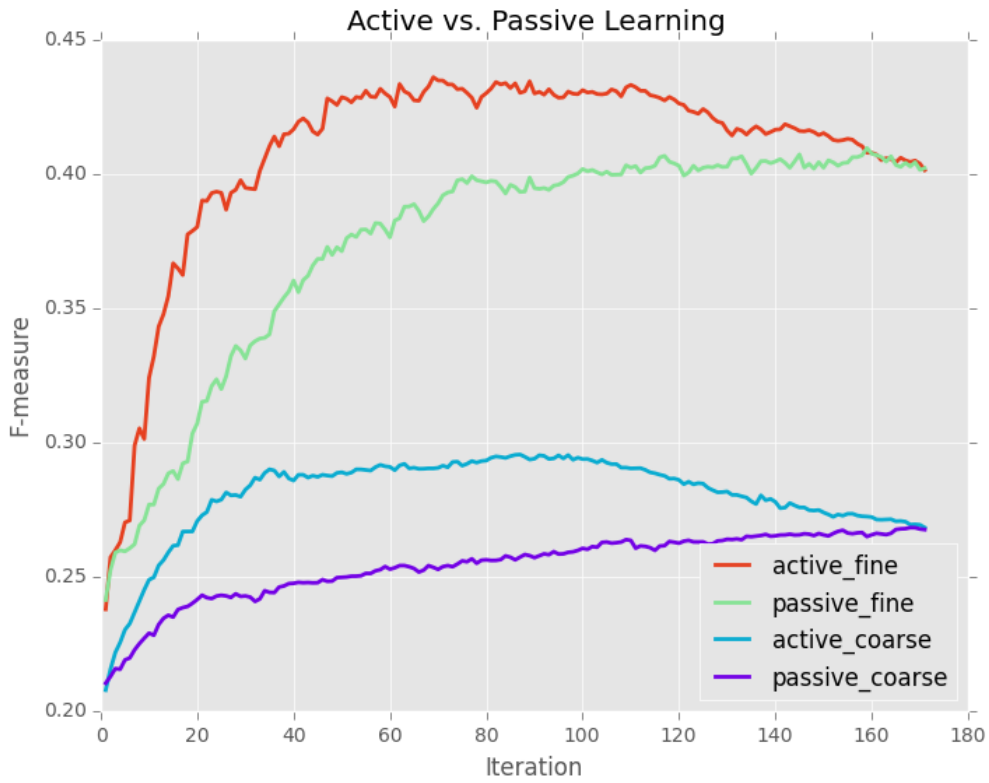

Figure: Both curves show a dominance of fine over coarse and Active over Passive.

Dynamically Adapting Purchase Proportions p or p'

HAL - Protein
James Duin
Introduction
Background
Exp. Setup
Conv. ML
Act. vs Pass.
HAL Results
BANDIT Results
Conclusions

- For round n, calculate gain g in terms of observed model change
- Calculate average round reward for each arm
- Calculate $\varepsilon_n = \min\left\{1, \frac{2}{n}\right\}$
- With probability $1 - \varepsilon_n$ play arm with highest current average reward for round n, otherwise explore
- After playing arm, run HAL with chosen p or p'

ARM STAY	ARM SWITCH
$r(n) = 0$	$r(n) = \begin{cases} -g(n)/\lvert g(n)\rvert & \text{if } p \to p' \\ g(n)/\lvert g(n)\rvert & \text{if } p' \to p \\ 0 & \text{if } p \to p \text{ or } p' \to p' \end{cases}$

Bibliography

HAL - Protein
James Duin
Introduction
Background
Exp. Setup
Conv. ML
Act. vs Pass.
HAL Results
BANDIT Results
Conclusions

- Y. Mo, S. D. Scott, and D. Downey, Learning hierarchically decomposable concepts with active over-labeling, in 2016 IEEE 16th International Conference on Data Mining (ICDM), Dec 2016, pp. 340349.
- J. Z. Juan Cui, Kevin Chiang, Prediction of nuclear and locally encoded mitochondrion. Lincoln, NE: Nebraska Gateway to Nutrigenomics 6th Annual Retreat, June 9 2014. [Online]. Available: http://cehs.unl.edu/nutrigenomics/nebraska-gateway-nutrigenomics-6th-annual-retreat/
- T. M. Mitchell, Machine Learning, 1st ed. New York, NY, USA: McGraw-Hill, Inc., 1997.

Bibliography

HAL - Protein

James Duin

Introduction

Background

Exp. Setup

Conv. ML

Act. vs Pass.

HAL Results

BANDIT Results

Conclusions

- L. Buitinck, G. Louppe, M. Blondel, F. Pedregosa, A. Mueller, O. Grisel, V. Niculae, P. Prettenhofer, A. Gramfort, J. Grobler, R. Layton, J. VanderPlas, A. Joly, B. Holt, and G. Varoquaux, API design for machine learning software: experiences from the scikit-learn project, in ECML PKDD Workshop: Languages for Data Mining and Machine Learning, 2013, pp. 108122
- D. Cotter, P. Guda, E. Fahy, and S. Subramaniam, Mitoproteome: mitochondrial protein sequence database and annotation system, Nucleic Acids Research, vol. 32, no. suppl1, p. D463, 2004. [Online]. Available: +http://dx.doi.org/10.1093/nar/gkh048

Bibliography

HAL - Protein
James Duin
Introduction
Background
Exp. Setup
Conv. ML
Act. vs Pass.
HAL Results
BANDIT Results
Conclusions

- J. Cui, L. Y. Han, H. Li, C. Y. Ung, Z. Q. Tang, C. J. Zheng, Z. W. Cao, and Y. Z. Chen, Computer prediction of allergen proteins from sequence-derived protein structural and physicochemical properties, Molecular Immunol-
ogy, vol. 44, no. 4, pp. 514 520, 2007. [Online]. Available: http://www.sciencedirect.com/science/article/pii/S016158900
- A. McCallum, R. Rosenfeld, T. M. Mitchell, and A. Y. Ng, Improving text classification by shrinkage in a hierarchy of classes, in Proceedings of the Fifteenth International Conference on Machine Learning, ser. ICML 98. San Francisco, CA, USA: Morgan Kaufmann Publishers Inc., 1998, pp. 359367. [Online]. Available: http://dl.acm.org/citation.cfm?id=645527.657461

Bibliography

HAL - Protein
James Duin
Introduction
Background
Exp. Setup
Conv. ML
Act. vs Pass.
HAL Results
BANDIT Results
Conclusions

- W. Jiang and Z. W. Ras, Multi-label automatic indexing of music by cascade classifiers, Web Intelli. and Agent Sys., vol. 11, no. 2, pp. 149170, Apr. 2013. [Online]. Available: http://dl.acm.org/citation.cfm?id=2590084.2590088
- etc.